


Where Stash sits in the AI-memory market — and where I fit

Farzan Aman Khan · github.com/cheese-cakee · July 2026

Supermemory and Mem0 sell memory technology to developers — invisible plumbing inside someone else's app. **Stash gives whole teams a shared brain**: one knowledge base that people and AI agents read and write together. That is a different market, and the more valuable one.

The market at a glance



Stash has the top-right quadrant to itself: team-wide knowledge, in a form humans can actually read.

Supermemory and Mem0 compete with each other — and with every AI lab that adds built-in memory to its models. Stash doesn't sit in that fight. It owns something none of them capture: **the working history of a team and its agents** — what was tried, what was decided, and why.

The comparison that matters

	STASH	SUPERMEMORY	MEMO
What it really is	A shared workspace for teams and their AI agents	A memory service developers plug into their apps	An open-source memory toolkit for developers
Who pays for it	Companies whose engineers work with AI agents	App builders	App builders
What the memory looks like	Ordinary files and pages — anyone can open, check, and fix them	Machine-managed data; hard to inspect or correct	Machine-managed data; hard to inspect or correct
Whose memory it is	The whole team's, shared	One user at a time	One user at a time
Can it run inside company walls?	Yes — fully self-hosted, open source	Cloud-first	Partly; managed platform for the full product

Three reasons the Stash approach wins

1 Agents already speak its language

Stash presents itself as a filesystem — the thing coding agents use all day, every day. No new interface to learn, no integration project. Adoption friction is close to zero.

2 Memory you can open and fix

When the other tools remember something wrong, nobody can see it — the error surfaces later, silently, in the AI's answers. In Stash a wrong memory is just a file: open it, correct it, done. For companies, that transparency is what makes AI memory trustworthy enough to buy.

3 It captures what nobody else does

Stash records what a team's agents attempted, decided, and abandoned. The most expensive knowledge in any engineering team is *why the last attempt failed* — and Stash is the only product of the three that keeps it.

Where I add value

Four things Stash will need as it grows — each matched to work I have already shipped.

WHAT STASH NEEDS

Keep secrets out of shared memory

Agent transcripts carry passwords and API keys. Shared, searchable memory turns one leak into everyone's leak.

Stay fast as the knowledge base grows

Search speed is the product. It has to hold up at 100× today's data.

Win the self-hosting customer

The companies that most need Stash are the ones that can't send data to a third-party cloud. Self-hosting has to feel effortless.

Security that enterprise buyers trust

Stash asks companies to store their most sensitive work product. Its own security posture is part of the pitch.

WHAT I'VE ALREADY DONE

I built exactly that scanner

Overtrust — my security tool that finds leaked secrets and credentials. I'd build the same protection into Stash's upload path.

Recent, proven performance work

12 merged performance and reliability PRs at Lucebox (an AI inference startup) in the past few weeks, on caching and search-ranking quality.

Infrastructure work at the Linux Foundation

LFX mentee (selected 1 of 151); built and shipped developer tooling end to end. I'd start with your open issue #563 on self-hosted setups.

Security fixes merged into major projects

Hardening work accepted into Ceph, CRIU, and Linux Foundation CI pipelines — including fixing real injection risks.

In short: I've read your code, I know your open issues, and the problems ahead of you are the ones I've already been working on. Give me trial work and I'll show you, not tell you.